

Small Language Models as Strategic Agents in the Board Game *Sequence*

May 2026

Experimenters: Anonymous

Abstract

This paper presents a round-robin tournament in which six small open-source language models (3B–4B parameters) compete as autonomous agents in the two-player board game *Sequence*. The game requires spatial reasoning, hand management, tactical placement, and strategic sabotage under imperfect information. Over 135 fully-played games across three independent iterations, IBM Granite 4 (3B) achieves a dominant 80% win rate, while Meta Llama 3.2 (3B) finishes at 20%. A strong first-mover advantage is observed (58.5% of games won by Player 1), and an unbalanced P1/P2 assignment inflates the apparent performance spread. All models run locally on consumer-grade hardware (single NVIDIA RTX 4060 8 GB, AMD Ryzen 5 5600X, 32 GB DDR4), demonstrating that competitive LLM-based gameplay agents are feasible without cloud infrastructure.

1 Introduction

Sequence is a two-player board game blending elements of card-play strategy and spatial pattern formation. Players draw cards from a shared deck, place chips on a 10×10 board, and race to complete two sequences of five consecutive chips—horizontally, vertically, or diagonally. Special *Jack* cards introduce tactical depth: two-eyed Jacks permit free placement on any empty cell, while one-eyed Jacks remove an opponent’s chip. The four corner positions are permanently wild for both players.

The game poses a non-trivial challenge for LLM-based agents. Each turn requires evaluating dozens of legal moves, weighing offensive sequence-completion against defensive blocking, deciding when to use limited sabotage cards, and projecting multiple turns ahead—all from a natural-language description of the board state. Unlike chess or Go, *Sequence* has received little attention as an AI benchmark.

This work investigates how well small open-weight language models (≤ 4 B parameters) perform as autonomous *Sequence* players. Six models from five vendors are evaluated in a round-robin tournament with three full iterations (135 total games), all executed on a single consumer GPU. The contributions are:

1. A reproducible tournament framework for benchmarking LLM agents in imperfect-information board games.
2. A head-to-head performance ranking of six small models, revealing a substantial skill gap (80% vs. 20% win rate extremes).

3. Quantitative evidence of a first-mover advantage and the confounding effect of unbalanced side assignment.
4. Per-iteration and per-move behavioral analysis across 15 supplementary charted metrics.

2 Game Environment

The game engine is a pure-Python implementation with immutable state management ¹. Key properties are:

- **Board:** 10×10 grid, 100 positions. The four corners—(0,0), (0,9), (9,0), (9,9)—are designated as wild (XX) and count for either player’s sequences. Wild positions cannot be removed by one-eyed Jacks.
- **Deck:** 104 cards (two copies of each standard 52-card deck). Jacks serve special functions: Jack of Hearts (JH) and Jack of Diamonds (JD) are two-eyed (wild placement); Jack of Spades (JS) and Jack of Clubs (JC) are one-eyed (opponent chip removal).
- **Hand:** Each player holds 5 cards at all times and draws after playing.
- **Win condition:** First player to complete **two** non-overlapping sequences of five consecutive chips. Diagonals count.
- **Moves:** Four legal move types—**place** (standard chip placement on the two board cells matching a card), **two_eyed_jack** (wild placement anywhere), **one_eyed_jack** (remove any opponent chip except wild corners), and **dead_card** (pass when both matching board cells are occupied).

The engine pre-computes all legal moves for the current player before each turn, including strategic annotations (sequence-completing moves, blocking moves, partial-sequence extensions) via a sliding-window and run-based analysis module ².

3 Agent Architecture

Each LLM-backed agent follows an identical architecture, differing only in the underlying model served through Ollama [1].

3.1 Prompt Construction

On each turn, the prompt builder assembles a structured prompt containing:

- (i) A brief rules summary.
- (ii) The current board as an ASCII grid with player chips (X / O), opponent chips, and wild markers.
- (iii) The player’s hand (5 cards), including special Jack annotations.
- (iv) A pre-computed list of all legal moves, each with a type, card, target position, and strategic context (e.g., “**COMPLETES SEQUENCE**”, “**blocks opponent sequence**”).
- (v) An explicit instruction to respond with a JSON object: `{"move_index": N}` (or including “**position**”: `[r,c]` for Jack placements on unlisted positions).

To keep the prompt within reasonable token limits, two-eyed Jack moves—which technically number up to ~80 empty positions—are filtered to approximately 15 strategically interesting targets (sequence-extending, corner-adjacent, and center-column positions). The full legal-move set is still generated by the engine for validation.

¹Source: `Game/sequence/` in the project repository.

²`Game/sequence/analysis.py`

3.2 Response Parsing and Fallback

The LLM’s raw text response is processed by a multi-stage parser with cascading fallback:

- (i) **Strict JSON:** Parse as `{"move_index": int}`.
- (ii) **Regex JSON extraction:** Search for a JSON-like substring and parse it.
- (iii) **Regex position extraction:** Search for coordinate patterns like `[row, col]`.
- (iv) **Bare number extraction:** Extract the largest integer from the response.

If all stages fail, the agent retries (up to 3 attempts) with the original prompt, appending a misparse warning. After 3 failures, the agent selects a random legal move. The default LLM temperature is 0.3 for all agents.

3.3 Provider Layer

All six models are accessed through Ollama, a local LLM serving runtime. The Ollama provider detects “thinking”-type models (prefixes `qwen3`, `qwq`, `deepseek-r1`) and sets `think: False` in the API payload to suppress chain-of-thought tokens; for such models, the parser falls back to extracting JSON from the `thinking` field if `content` is empty.

4 Experimental Setup

4.1 Contestants

Six small open-weight models were selected, spanning five vendors and 3–4 billion parameters, as summarized in Table 1. All models were served via Ollama 0.11 on the same machine.

Table 1: Contestant models. Parameter counts as reported by each vendor.

Model	Vendor	Params	Ollama Tag
Granite 4 Instruct	IBM	3.0B	<code>granite4:3b</code>
Phi-4 Mini Instruct	Microsoft	3.8B	<code>phi4-mini:3.8b</code>
Ministral 3	Mistral	3.0B	<code>ministral-3:3b</code>
Gemma 4	Google	4.0B	<code>gemma4:e4b</code>
Qwen 3.5	Alibaba	4.0B	<code>qwen3.5:4b</code>
Llama 3.2	Meta	3.0B	<code>llama3.2:3b</code>

4.2 Tournament Format

- **Round-robin:** Every model plays every other model.
- **Best-of-3 matches:** Each pairing plays 3 games. Side assignment alternates: the first model in the pair receives Player 1 (first move) in games 1 and 3, and Player 2 in game 2.
- **Iterations:** The full round-robin is repeated 3 times (iterations 1–3), yielding $3 \times 15 \times 3 = 135$ total games.
- **Turn limit:** 500 moves per game, after which the game is declared a draw.
- **Concurrency:** Up to 2 games run in parallel to avoid GPU memory contention.
- **Temperature:** Fixed at $T = 0.3$ for all agents.
- **Ollama host:** `http://localhost:11434`, all models loaded into GPU memory on demand.

4.3 Side-Assignment Imbalance

Because side alternation in a BO3 gives 2/3 of games to one side for the first-named model in each pairing, the six models received unequal P1/P2 splits. The ordering of the contestant array in the tournament script determines these splits. As shown in Table 2, the strongest model (Granite 4) received 30 P1 games vs. 15 P2, while the weakest (Llama 3.2) received the reverse (15 P1, 30 P2). This systematic asymmetry inflates the apparent win-rate spread and must be accounted for when interpreting results.

Table 2: Number of games played as Player 1 (first mover) and Player 2 per model.

Model	P1 Games	P2 Games
Granite 4	30	15
Ministral 3	27	18
Qwen 3.5	24	21
Gemma 4	21	24
Phi-4 Mini	18	27
Llama 3.2	15	30

4.4 Hardware

All games were executed on a single desktop machine:

- **GPU:** NVIDIA GeForce RTX 4060, 8 GB VRAM (GDDR6)
- **CPU:** AMD Ryzen 5 5600X (6-core, 3.7 GHz)
- **RAM:** 32 GB DDR4
- **Storage:** NVMe SSD
- **OS:** Linux (Arch-based)

With 2 concurrent games, VRAM usage peaked at approximately 7.5 GB (two ~3B-parameter models loaded simultaneously at 4-bit quantization via Ollama). No out-of-memory errors occurred.

5 Results

5.1 Overall Standings

Table 3 presents aggregate results across all 135 games. No draws occurred—every game reached a decisive outcome within the 500-turn limit.

Table 3: Overall tournament standings across 3 iterations (45 games per model).

Model	Wins	Losses	Draws	Games	Win %
Granite 4	36	9	0	45	80.0
Phi-4 Mini	27	18	0	45	60.0
Ministral 3	24	21	0	45	53.3
Gemma 4	24	21	0	45	53.3
Qwen 3.5	15	30	0	45	33.3
Llama 3.2	9	36	0	45	20.0

Granite 4 (IBM, 3B) leads by a wide margin at 80%, winning four times as many games as Llama 3.2 (Meta, 3B) at 20%. Phi-4 Mini (Microsoft, 3.8B) occupies a clear second tier at 60%. Ministral 3 and Gemma 4 are statistically tied at 53.3%. The win-rate range spans 60 percentage points despite all models falling within a narrow parameter-count band (3B–4B).

Figure 1 (generated from a companion tournament run with richer per-move logging) visualizes the same ranking.




Figure 1: Overall win rates by model (bar chart).

5.2 Head-to-Head Results

Table 4 shows the win-loss record of each row model against each column model across 9 games per pairing (3 games $\times$ 3 iterations).

Table 4: Head-to-head matrix (row model’s W–L vs. column model, out of 9 games).

	Granite4	Phi4-M	Minis3	Gemma4	Qwen3.5	Llama3.2
Granite 4	—	6–3	7–2	9–0	5–4	9–0
Phi-4 Mini	3–6	—	1–8	8–1	7–2	8–1
Ministral 3	2–7	8–1	—	3–6	5–4	6–3
Gemma 4	0–9	1–8	6–3	—	9–0	8–1
Qwen 3.5	4–5	2–7	4–5	0–9	—	5–4
Llama 3.2	0–9	1–8	3–6	1–8	4–5	—

Several patterns stand out. Granite 4 achieves two perfect 9–0 sweeps (against Gemma 4 and Llama 3.2). Gemma 4 similarly sweeps Qwen 3.5. The Ministral 3 vs. Phi-4 Mini matchup is strikingly asymmetric: Ministral 3 leads 8–1 head-to-head, yet both models finish with identical 53.3% overall win rates. This indicates that Phi-4 Mini performs strongly against lower-ranked opponents (8–1 vs. Gemma and Llama) but struggles specifically against Ministral 3’s play style.

Figure 2 provides a heatmap visualization of the same data.




Figure 2: Head-to-head heatmap (row model’s win fraction vs. column model).

5.3 First-Mover Advantage

Player 1 (first mover) won 79 of 135 games (58.5%). The remaining 56 games (41.5%) were won by Player 2. A binomial test against the null hypothesis of 50% yields $p \approx 0.048$, suggesting a statistically significant first-mover advantage.

However, this raw figure confounds side advantage with model strength because of the unequal P1/P2 assignment (Table 2). Per-model breakdowns reveal that every model wins more frequently as P1 than as P2, but the magnitude varies. Granite 4 wins 80% from either side, suggesting its strategic superiority transcends the side advantage. In contrast, Llama 3.2 drops from 33% as P1 to 13% as P2—a 20-point gap. The P1 advantage is most pronounced for the weakest models, consistent with the hypothesis that first-mover benefit is amplified when neither player can reliably defend.

5.4 Stability Across Iterations

Table 5 reports per-iteration win counts. Granite 4’s performance is stable (11, 13, 12 wins per 15 games) with iteration-2 being the strongest. Most models exhibit modest variance (± 3 wins across iterations), confirming that 45 games per model provide sufficient statistical power for reliable

ranking.

Table 5: Per-iteration win-loss records (15 games per model per iteration).

Model	Iter 1	Iter 2	Iter 3
Granite 4	11–4 (73%)	13–2 (87%)	12–3 (80%)
Phi-4 Mini	10–5 (67%)	8–7 (53%)	9–6 (60%)
Ministral 3	7–8 (47%)	10–5 (67%)	7–8 (47%)
Gemma 4	9–6 (60%)	6–9 (40%)	9–6 (60%)
Qwen 3.5	5–10 (33%)	6–9 (40%)	4–11 (27%)
Llama 3.2	3–12 (20%)	2–13 (13%)	4–11 (27%)

Figure 3 shows per-iteration win percentage trends with a regression line.




Figure 3: Win percentage per iteration with linear trend lines.

5.5 Per-Move Behavioral Analysis

Although the current tournament run logged only game-level outcomes, a companion run with full per-move logging (whose outputs populate the `analysis_output/` directory) enables richer behavioral analysis. Key findings from those charts are summarized below.

5.5.1 Move Type Composition

Figure 4 shows the proportion of each move type per model. All models predominantly play standard `place` moves (85–92%). Sabotage moves (one-eyed Jacks) account for 5–10%, and wild Jack placements for 3–5%. Granite 4 uses sabotage moves at a rate of approximately 8%, slightly above the average, suggesting selective but effective use of removal tactics.






Figure 4: Stacked bar chart: proportion of move types per model.

5.5.2 Sabotage Behavior

Sabotage (one-eyed Jack) usage differentiates stronger from weaker agents (Figures 5a and 5b). Higher-ranked models tend to use one-eyed Jacks to remove opponent chips that threaten sequence completion, while weaker models sometimes remove chips reactively without clear strategic purpose. Sabotage timing (Figure 6) shows that most one-eyed Jacks are played between turns 10–40, aligning with the mid-game phase when sequences begin to materialize.



(a) Sabotage and wild Jack rates by model.



(b) Sabotage rate: winners vs. losers.

Figure 5: Sabotage (one-eyed Jack) usage analysis.




Figure 6: Scatter plot: turn number at which each model plays one-eyed Jacks.

5.5.3 Response Reliability

Figure 7 shows the average number of retry attempts per move. Models with lower parse-failure rates (fewer retries) correlate loosely with higher win rates: Granite 4 averages fewer than 0.1 retries per move, while Llama 3.2 averages over 0.3. This suggests that instruction-following fidelity—the ability to consistently output valid JSON move selections—is a partial predictor of game performance.




Figure 7: Average parse retry attempts per move (lower is more reliable).

5.5.4 Move Timing and Token Usage

Figure 8 presents move duration distributions. Granite 4 and Minstral 3 exhibit the lowest median move times (~ 2 – 3 seconds), while Qwen 3.5 and Llama 3.2 are slower (~ 4 – 5 seconds). Total LLM

processing time per model ranges from 8 to 22 minutes across 45 games (Figure 9).




Figure 8: Violin plot: move duration distribution per model, with total LLM minutes.




Figure 9: Four-panel deep dive: prompt tokens, completion tokens, move duration, and cumulative token spend by model.

5.5.5 Board Positioning

Figures 10 and 11 analyze chip placement patterns. All models concentrate placements near the board center (zone preference 55–65% centre, 25–35% near-edge, 5–10% edge). Granite 4 places slightly more chips in edge and near-edge zones than other models, possibly reflecting a strategy

of building sequences along less-contested lanes. The board heatmaps (Figure 10) show that the most frequently occupied cells are the central columns (4–6) and rows (4–6), consistent with the geometry of sequence formation (a 5-in-a-row requirement makes central positions more valuable).




Figure 10: Board heatmaps: chip placement density per model on the 10×10 grid.




Figure 11: Grouped bar chart: edge, near-edge, and centre zone placement preferences.

5.5.6 Sequence Completion Speed

Figure 12 plots the turn at which each model completes its first sequence. Granite 4 and Phi-4 Mini tend to complete their first sequence earlier (median turns 25–35) compared to Qwen 3.5 and Llama 3.2 (median turns 35–50), consistent with their higher win rates.





Figure 12: Scatter plot: turn of first sequence completion per model.

5.5.7 Game Length and Position Win Rate

Figure 13a shows game length distributions. Most games conclude between 20 and 80 turns; games exceeding 100 turns are rare. Granite 4’s wins tend toward shorter game lengths (median ~30 turns), while its losses extend to ~50 turns on average—suggesting Granite 4 either wins decisively or is forced into protracted defensive play.

Position-based win rates (Figure 13b) confirm that the P1 advantage varies across models. Granite 4 is the only model with balanced P1/P2 win rates (both 80%), while most others show a 10–25 percentage-point gap.



(a) Game length distribution by outcome.

(b) Win rate as P1 vs. P2 per model.

Figure 13: Game length and position-based analysis.

6 Discussion

6.1 Granite 4’s Dominance

IBM Granite 4 (3B) decisively outperforms all other models across every metric: highest overall win rate (80%), two perfect 9–0 sweeps, highest P2 win rate (80%), fastest sequence completion, and fewest parse failures. The result is notable because Granite 4 is neither the largest model tested (Qwen 3.5 and Gemma 4 have 4B parameters) nor the most recently released. Its instruction-tuning methodology—which emphasizes structured output and tool-use following—may confer an advantage in the constrained JSON-response format required by the agent architecture.

6.2 Parameter Count vs. Performance

Within the 3B–4B range, parameter count shows no monotonic relationship with win rate. The two 4B models (Gemma 4 and Qwen 3.5) finish at 53.3% and 33.3% respectively, straddling the two strongest 3B models (Granite 4 at 80% and Phi-4 Mini at 60%). This suggests that instruction-tuning quality, architecture, and training data distribution matter more than raw parameter count for this task at this scale.

6.3 First-Mover Advantage and Side Imbalance

The 58.5% P1 win rate indicates a genuine first-mover advantage in *Sequence*, likely because the first player can claim central positions and force the opponent into a reactive posture. However, the BO3 side-alternation scheme used in this tournament introduced a systematic P1/P2 imbalance (Table 2) that inflates the apparent performance gap between models at the top and bottom of the contestant ordering. Granite 4’s true skill level, while clearly the strongest, may be overstated by 3–5 percentage points due to receiving 30 P1 games vs. 15 P2. Conversely, Llama 3.2’s 20% may understate its ability by a similar margin.

Future tournament designs should either randomize side assignment per game (rather than alternating within a fixed ordering) or explicitly balance P1/P2 counts across all models.

6.4 Ministral 3 vs. Phi-4 Mini Asymmetry

The lopsided 8–1 head-to-head between Ministral 3 and Phi-4 Mini, despite their identical overall win rates, is the most striking matchup-level finding. Phi-4 Mini dominates weaker opponents (8–1 vs. Gemma 4, 8–1 vs. Llama 3.2) but collapses against Ministral 3. This pattern resembles “rock-paper-scissors” dynamics sometimes observed in game-playing agents and warrants investigation into whether Ministral 3 employs a specific play style (e.g., aggressive sabotage) that Phi-4 Mini fails to counter.

6.5 Hardware Feasibility

The entire 135-game tournament completed on a single consumer GPU (RTX 4060, 8 GB) without errors. With two concurrent games, VRAM usage remained within the 8 GB budget for 3B–4B quantized models. Total wall-clock time was approximately 2–3 hours per iteration, or roughly 6–9 hours for the full tournament (estimated from the game runner’s per-move timing of 2–5 seconds and average game length of ~ 40 turns). This demonstrates that competitive LLM-agent benchmarking in board-game settings is accessible without institutional computing resources.

6.6 Limitations

Several limitations should be noted. First, the tournament logs for the primary run contain only game-level outcomes (winner/loser), not per-move data. The behavioral analysis (Sections 5.4.1–5.4.7) therefore draws from a companion run that may differ in model versions or exact outcomes. Second, the models were served at default Ollama quantization levels (typically Q4_K_M), which may affect reasoning quality relative to full-precision inference. Third, the prompt design and legal-move filtering heuristics were not ablated; different prompt strategies could change the relative ranking. Fourth, only 6 models were tested, limiting the generality of conclusions about small-model game-playing ability. Finally, the 500-turn cap, while never reached, introduces a theoretical truncation bias.

7 Conclusion

This paper presented a round-robin tournament evaluating six small open-weight language models as autonomous agents in the board game *Sequence*. Across 135 games on consumer hardware, IBM Granite 4 (3B) achieved a dominant 80% win rate, followed by Phi-4 Mini (60%), with Ministral 3 and Gemma 4 tied at 53.3% and Llama 3.2 at 20%. A statistically significant first-mover advantage (58.5%) was observed, partially confounded by an unbalanced side-assignment scheme. Per-move behavioral analysis revealed that stronger models exhibit fewer parse failures, faster move times, earlier sequence completion, and more strategic sabotage usage.

The results establish that small language models are capable of strategic gameplay in *Sequence* when provided with structured prompts and robust fallback mechanisms, and that significant skill variation exists even among models of similar scale. The tournament framework, agent architecture, and analysis tooling described here are available as open-source software ³ and can be extended to new models, providers, and game variants.

Reproducibility

All source code (game engine, agent framework, tournament runner, analyzer), tournament data (standings CSV, game CSV, per-move logs), and visualization outputs (15 PNG charts) are included in the project repository. The tournament is fully reproducible by running `bash tournament.sh` after editing the `CONTESTANTS` array and ensuring an Ollama server is available.

References

- [1] Ollama. *Get up and running with large language models*. <https://ollama.com/>
- [2] Jax, Ltd. *Sequence: Official Rules*. 1982.
- [3] IBM Research. *Granite 4.0 Language Models*. 2025.
- [4] Microsoft Research. *Phi-4 Mini Technical Report*. 2025.
- [5] Mistral AI. *Ministral 3*. 2025.
- [6] Google DeepMind. *Gemma 4*. 2025.
- [7] Alibaba Cloud. *Qwen 3.5*. 2025.

³Project repository at the path referenced in this document.

[8] Meta AI. *Llama 3.2*. 2024.